

# <sup>1</sup> LabelFusion: Learning to Fuse LLMs and Transformer Classifiers for Robust Text Classification

<sup>3</sup> Michael Schlee<sup>1</sup>, Christoph Weisser<sup>1</sup>, Timo Kivimäki<sup>2</sup>, Melchy Mashiku<sup>4</sup>, and  
<sup>4</sup> Benjamin Saefken<sup>3</sup>

<sup>5</sup> 1 Centre for Statistics, Georg-August-Universität Göttingen, Germany <sup>2</sup> Department of Politics and  
<sup>6</sup> International Studies, University of Bath, Bath, UK <sup>3</sup> Institute of Mathematics, Clausthal University of  
<sup>7</sup> Technology, Clausthal-Zellerfeld, Germany <sup>4</sup>

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a

Creative Commons Attribution 4.0  
International License ([CC BY 4.0](#)).

## <sup>8</sup> Summary

<sup>9</sup> LabelFusion is a fusion ensemble for text classification that learns to combine a traditional  
<sup>10</sup> transformer-based classifier (e.g., RoBERTa) with one or more Large Language Models (LLMs)  
<sup>11</sup> such as OpenAI GPT, Google Gemini, or DeepSeek to deliver accurate and cost-aware predictions  
<sup>12</sup> across multi-class and multi-label tasks. The package provides a simple high-level interface  
<sup>13</sup> (AutoFusionClassifier) that trains the full pipeline end-to-end with minimal configuration,  
<sup>14</sup> and a flexible API for advanced users. Under the hood, LabelFusion concatenates vector signals  
<sup>15</sup> from the ML backbone (logits) and LLM(s) (per-class scores) and trains a compact multi-layer  
<sup>16</sup> perceptron (FusionMLP) to produce the final prediction. This learned fusion approach captures  
<sup>17</sup> complementary strengths of LLM reasoning and traditional transformer-based classifiers,  
<sup>18</sup> yielding robust performance across domains—achieving 92.4% accuracy on AG News topic  
<sup>19</sup> classification—while enabling practical trade-offs between accuracy, latency, and cost.

## Statement of Need

<sup>20</sup> Modern text classification spans diverse scenarios—from sentiment analysis to complex topic  
<sup>21</sup> tagging—often under constraints that vary per deployment (throughput, cost ceilings, data  
<sup>22</sup> privacy). While transformer classifiers such as BERT/RoBERTa achieve strong supervised  
<sup>23</sup> performance ([Devlin et al., 2018](#); [Liu et al., 2019](#)), frontier LLMs can excel in low-data,  
<sup>24</sup> ambiguous, or cross-domain settings ([OpenAI, 2023](#)). No single model family is typically  
<sup>25</sup> uniformly best: LLMs are powerful, but comparatively costly, whereas fine-tuned transformers  
<sup>26</sup> are efficient but may struggle with out-of-distribution cases.

<sup>27</sup> LabelFusion addresses this gap by: (1) exposing a minimal “AutoFusion” interface that trains a  
<sup>28</sup> learned combination of an ML backbone and one or more LLMs; (2) supporting both multi-class  
<sup>29</sup> and multi-label classification; (3) providing a lightweight fusion learner that directly fits on LLM  
<sup>30</sup> scores and ML logits; and (4) integrating cleanly with existing ensemble utilities. Researchers  
<sup>31</sup> and practitioners can therefore leverage LLMs where they add value while retaining the speed  
<sup>32</sup> and determinism of transformer models.

## <sup>34</sup> State of the Field

<sup>35</sup> In applied NLP, common tools such as scikit-learn ([Pedregosa et al., 2011](#)) and Hugging Face  
<sup>36</sup> Transformers ([Wolf et al., 2019](#)) offer strong baselines but do not provide a learned fusion of  
<sup>37</sup> LLMs with supervised transformers. Orchestration frameworks (e.g., LangChain) focus on tool  
<sup>38</sup> use rather than classification ensembles. LabelFusion contributes a focused, production-minded  
<sup>39</sup> implementation of a small learned combiner that operates on per-class signals from both model  
<sup>40</sup> families.

## 41    **Functionality and Design**

42    LabelFusion consists of three layers:

- 43    ▪ ML component: a RoBERTa-style classifier produces per-class logits for input texts.
- 44    ▪ LLM component(s): provider-specific classifiers (OpenAI, Gemini, DeepSeek) return per-class scores via prompting. Scores can be cached to minimize API calls when cache locations are provided.
- 45    ▪ Fusion component: a compact MLP concatenates ML logits and LLM scores and outputs fused logits. The ML backbone is trained/fine-tuned with a small learning rate; the fusion MLP uses a higher rate, enabling rapid adaptation without destabilizing the encoder.

50    Key features:

- 51    ▪ **Multi-class and multi-label support** with consistent data structures and unified training pipeline.
- 52    ▪ **Optional LLM response caching** reuses on-disk predictions when cache paths are supplied, with dataset-hash validation to guard against stale files.
- 53    ▪ **Batched scoring** processes multiple texts efficiently with configurable batch sizes for both ML tokenization and LLM API calls.
- 54    ▪ **Results management** via ResultsManager tracks experiments, stores predictions, computes metrics, and enables reproducible research workflows.
- 55    ▪ **Flexible interfaces**: Command-line training via `train_fusion.py` with YAML configs for research; or minimal AutoFusion API for quick deployment.
- 56    ▪ **Composable design**: LabelFusion can serve as a strong base learner in higher-level ensembles (e.g., voting/weighted combinations of multiple fusion models).

## 63    **Minimal Example (AutoFusion)**

```
from textclassify import AutoFusionClassifier

config = {
    'llm_provider': 'deepseek',
    'label_columns': ['positive', 'negative', 'neutral']
}

clf = AutoFusionClassifier(config)
clf.fit(train_dataframe)           # trains ML backbone, gathers LLM scores, fits fusi
pred = clf.predict(["This is amazing!"]) # fused prediction
```

## 64    **CLI and Configuration**

65    Users can generate a starter config and train via the command line:

- 66    ▪ Create config: `python train_fusion.py --create-config fusion_config.yaml`
- 67    ▪ Train: `python train_fusion.py --config fusion_config.yaml`
- 68    ▪ Optional test data and output artifacts are also supported.

## 69    **Quality Control**

70    The repository ships legacy unit tests under `tests/evaluation/old/` that cover configuration handling, core types, and package integration. Fusion-specific logic is currently exercised through CLI-driven workflows and notebooks that run end-to-end training with deterministic seeds where applicable.

74    Evaluation scripts (`tests/evaluation/`) provide comprehensive benchmarking on standard datasets:  
 - **AG News** ([Zhang et al., 2015](#)): 4-class topic classification with experiments across varying training data sizes (20%–100%)  
 - **GoEmotions** ([Demszky et al., 2020](#)): 28-class multi-label emotion classification for validating multi-label fusion performance

<sup>78</sup> LLM scoring paths implement retries and disk caching; transformer training supports standard  
<sup>79</sup> sanity checks (overfit a small batch, reduced batch sizes for constrained hardware). Metrics  
<sup>80</sup> (accuracy/F1, per-label scores) are computed automatically and stored with run artifacts to  
<sup>81</sup> facilitate regression tracking and reproducibility.

## <sup>82</sup> Availability and Installation

<sup>83</sup> LabelFusion is distributed as part of the `textclassify` package under the MIT license and  
<sup>84</sup> is available at <https://github.com/DataandAIResearch/LabelFusion>. The fusion components  
<sup>85</sup> require Python 3.8+ and common scientific Python dependencies (PyTorch, transformers,  
<sup>86</sup> scikit-learn, numpy, pandas, PyYAML). Optional plotting depends on matplotlib/seaborn.  
<sup>87</sup> Installation and quick-start snippets are provided in the README and FUSION\_README.md.

## <sup>88</sup> Production-Ready Features

<sup>89</sup> Beyond the core fusion methodology, LabelFusion includes features for practical deployment:

- <sup>90</sup> **LLM Response Caching:** Optional disk-backed caches reuse prior predictions when cache  
<sup>91</sup> paths are supplied, with dataset hashes to flag inconsistent inputs.
- <sup>92</sup> **Results Management:** Built-in `ResultsManager` tracks experiments, stores predictions,  
<sup>93</sup> and computes metrics automatically. Supports comparison across runs and configuration  
<sup>94</sup> tracking.
- <sup>95</sup> **Batch Processing:** Efficient batched scoring of texts with configurable batch sizes for  
<sup>96</sup> both ML and LLM components.

## <sup>97</sup> Impact and Use Cases

### <sup>98</sup> Empirical Performance

<sup>99</sup> LabelFusion has been evaluated on standard benchmark datasets to validate its effectiveness.  
<sup>100</sup> Key findings demonstrate consistent improvements over individual model components:

#### <sup>101</sup> AG News Topic Classification

<sup>102</sup> Evaluation on the AG News dataset ([Zhang et al., 2015](#)) (4-class topic classification) with  
<sup>103</sup> 5,000 test samples shows:

Training Data	Model	Accuracy	F1-Score	Precision	Recall
20% (800)	<b>Fusion</b>	<b>92.2%</b>	<b>0.922</b>	0.923	0.922
20% (800)	RoBERTa	89.8%	0.899	0.902	0.898
20% (800)	OpenAI	84.4%	0.844	0.857	0.844
40% (1,600)	<b>Fusion</b>	<b>92.2%</b>	<b>0.922</b>	0.924	0.922
40% (1,600)	RoBERTa	91.0%	0.911	0.913	0.910
40% (1,600)	OpenAI	84.4%	0.844	0.857	0.844
100% (4,000)	<b>Fusion</b>	<b>92.4%</b>	<b>0.924</b>	0.926	0.924
100% (4,000)	RoBERTa	92.2%	0.922	0.923	0.922
100% (4,000)	OpenAI	84.4%	0.844	0.857	0.844

<sup>104</sup> **Key Observations:** - Fusion consistently outperforms individual models across all training  
<sup>105</sup> data sizes - With only 20% training data, Fusion achieves 92.2% accuracy—matching its  
<sup>106</sup> performance with full data - Demonstrates superior **data efficiency**: fusion learning extracts  
<sup>107</sup> maximum value from limited examples - RoBERTa alone requires 100% of data to approach  
<sup>108</sup> Fusion's 20% performance - LLM (OpenAI) shows stable but lower performance, highlighting  
<sup>109</sup> the value of combining approaches

110 These results validate that learned fusion captures complementary strengths: the LLM provides  
111 robust reasoning even with limited training data, while the ML backbone adds efficiency and  
112 domain-specific patterns.

### 113 Application Domains

114 Learned fusion excels in scenarios where model strengths complement each other:

- 115   ■ **Customer feedback analysis** with nuanced multi-label taxonomies where LLMs handle  
116 ambiguous sentiment while ML models efficiently process clear cases
- 117   ■ **Content moderation** where uncertain cases benefit from LLM reasoning while rou-  
118 tine items rely on the fast ML backbone, enabling real-time processing with accuracy  
119 guarantees
- 120   ■ **Scientific literature classification** across heterogeneous topics where domain shift is  
121 common and LLMs provide robustness to new terminology
- 122   ■ **Low-resource settings** where limited training data is available but task complexity requires  
123 sophisticated reasoning

124 The approach enables pragmatic cost control (e.g., the fusion layer learns when to rely more  
125 heavily on the efficient ML backbone versus the more expensive LLM signal) while retaining a  
126 single trainable decision surface that optimizes for the specific deployment constraints.

### 127 Acknowledgements

128 We thank contributors and users who reported issues and shared datasets. LabelFusion  
129 builds on the open-source ecosystem, notably Hugging Face Transformers (Wolf et al., 2019),  
130 scikit-learn (Pedregosa et al., 2011), PyTorch (Paszke et al., 2019), and LLM provider SDKs.  
131 We acknowledge the use of the AG News and GoEmotions benchmark datasets for evaluation.

### 132 References

- 133 Demszky, D., Movshovitz-Attias, D., Ko, J., Cowen, A., Nemade, G., & Ravi, S. (2020).  
134 GoEmotions: A dataset of fine-grained emotions. *Proceedings of the 58th Annual Meeting*  
135 *of the Association for Computational Linguistics*, 4040–4054.
- 136 Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of deep  
137 bidirectional transformers for language understanding. *arXiv Preprint arXiv:1810.04805*.
- 138 Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer,  
139 L., & Stoyanov, V. (2019). RoBERTa: A robustly optimized BERT pretraining approach.  
140 *arXiv Preprint arXiv:1907.11692*.
- 141 OpenAI. (2023). GPT-4 technical report. *arXiv Preprint arXiv:2303.08774*.
- 142 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin,  
143 Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-  
144 performance deep learning library. *Advances in Neural Information Processing Systems*,  
145 32.
- 146 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M.,  
147 Prettenhofer, P., Weiss, R., Dubourg, V., & others. (2011). Scikit-learn: Machine learning  
148 in python. *Journal of Machine Learning Research*, 12(Oct), 2825–2830.
- 149 Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T.,  
150 Louf, R., Funtowicz, M., & others. (2019). HuggingFace’s transformers: State-of-the-art  
151 natural language processing. *arXiv Preprint arXiv:1910.03771*.
- 152 Zhang, X., Zhao, J., & LeCun, Y. (2015). Character-level convolutional networks for text  
153 classification. *Advances in Neural Information Processing Systems*, 28, 649–657.