

¹ LabelFusion: Learning to Fuse LLMs and Transformer Classifiers for Robust Text Classification

³ Michael Schlee¹, Christoph Weisser¹, Timo Kivimäki², Melchizedek
⁴ Mashiku⁴, and Benjamin Saefken³

⁵ 1 Centre for Statistics, Georg-August-Universität Göttingen, Germany ² Department of Politics and
⁶ International Studies, University of Bath, Bath, UK ³ Institute of Mathematics, Clausthal University of
⁷ Technology, Clausthal-Zellerfeld, Germany ⁴ Tanaq Management Services LLC, Contracting Agency to
⁸ the Division of Viral Diseases, Centers for Disease Control and Prevention, Chamblee, Georgia, USA

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).
21
22
23

⁹ Summary

¹⁰ LabelFusion is a fusion ensemble for text classification that learns to combine a traditional
¹¹ transformer-based classifier (e.g., RoBERTa) with one or more Large Language Models (LLMs)
¹² such as OpenAI GPT, Google Gemini, or DeepSeek to deliver accurate and cost-aware predictions
¹³ across multi-class and multi-label tasks. The package provides a simple high-level interface
¹⁴ (AutoFusionClassifier) that trains the full pipeline end-to-end with minimal configuration,
¹⁵ and a flexible API for advanced users. Under the hood, LabelFusion concatenates vector signals
¹⁶ from the ML backbone (logits) and LLM(s) (per-class scores) and trains a compact multi-layer
¹⁷ perceptron (FusionMLP) to produce the final prediction. This learned fusion approach captures
¹⁸ complementary strengths of LLM reasoning and traditional transformer-based classifiers,
¹⁹ yielding robust performance across domains—achieving 92.4% accuracy on AG News topic
²⁰ classification—while enabling practical trade-offs between accuracy, latency, and cost.

²¹ Statement of Need

²² Modern text classification spans diverse scenarios, from sentiment analysis (Kant et al., 2024;
Luber et al., 2021; Thormann et al., 2021) to complex topic tagging (Kant et al., 2022; A.
Thielmann, Weisser, Krenz, & Säfken, 2021; A. Thielmann, Weisser, & Krenz, 2021; A. F.
Thielmann et al., 2024), often under constraints that vary per deployment (throughput, cost
ceilings, data privacy). While transformer classifiers such as BERT/RoBERTa achieve strong
supervised performance (Devlin et al., 2018; Liu et al., 2019), frontier LLMs can excel in
low-data, ambiguous, or cross-domain settings (OpenAI, 2023). No single model family is
typically uniformly best: LLMs are powerful, but comparatively costly, whereas fine-tuned
transformers are efficient but may struggle with out-of-distribution cases.

³¹ LabelFusion addresses this gap by: (1) exposing a minimal “AutoFusion” interface that trains a
³² learned combination of an ML backbone and one or more LLMs; (2) supporting both multi-class
³³ and multi-label classification; (3) providing a lightweight fusion learner that directly fits on LLM
³⁴ scores and ML logits; and (4) integrating cleanly with existing ensemble utilities. Researchers
³⁵ and practitioners can therefore leverage LLMs where they add value while retaining the speed
³⁶ and determinism of transformer models.

³⁷ State of the Field

³⁸ In applied NLP, common tools such as scikit-learn (Pedregosa et al., 2011) and Hugging Face
³⁹ Transformers (Wolf et al., 2019) offer strong baselines but do not provide a learned fusion of
⁴⁰ LLMs with supervised transformers. Orchestration frameworks (e.g., LangChain) focus on tool
⁴¹ use rather than classification ensembles. LabelFusion contributes a focused, production-minded

⁴² implementation of a small learned combiner that operates on per-class signals from both model
⁴³ families.

⁴⁴ Functionality and Design

⁴⁵ LabelFusion consists of three layers:

- ⁴⁶ ▪ ML component: a RoBERTa-style classifier produces per-class logits for input texts.
- ⁴⁷ ▪ LLM component(s): provider-specific classifiers (OpenAI, Gemini, DeepSeek) return
- ⁴⁸ per-class scores via prompting. Scores can be cached to minimize API calls when cache
- ⁴⁹ locations are provided.
- ⁵⁰ ▪ Fusion component: a compact MLP concatenates ML logits and LLM scores and outputs
- ⁵¹ fused logits. The ML backbone is trained/fine-tuned with a small learning rate; the fusion
- ⁵² MLP uses a higher rate, enabling rapid adaptation without destabilizing the encoder.

⁵³ Key features:

- ⁵⁴ ▪ **Multi-class and multi-label support** with consistent data structures and unified training
- ⁵⁵ pipeline.
- ⁵⁶ ▪ **Optional LLM response caching** reuses on-disk predictions when cache paths are supplied,
- ⁵⁷ with dataset-hash validation to guard against stale files.
- ⁵⁸ ▪ **Batched scoring** processes multiple texts efficiently with configurable batch sizes for both
- ⁵⁹ ML tokenization and LLM API calls.
- ⁶⁰ ▪ **Results management** via ResultsManager tracks experiments, stores predictions, com-
- ⁶¹ putes metrics, and enables reproducible research workflows.
- ⁶² ▪ **Flexible interfaces**: Command-line training via `train_fusion.py` with YAML configs for
- ⁶³ research; or minimal AutoFusion API for quick deployment.
- ⁶⁴ ▪ **Composable design**: LabelFusion can serve as a strong base learner in higher-level
- ⁶⁵ ensembles (e.g., voting/weighted combinations of multiple fusion models).

⁶⁶ We support both multi-class setups (one label per input) and multi-label scenarios (multi-

⁶⁷ ple labels per input), and point readers to Appendix A for formal definitions and training

⁶⁸ implications.

⁶⁹ Minimal Example (AutoFusion)

```
from textclassify import AutoFusionClassifier

config = {
    'llm_provider': 'deepseek',
    'label_columns': ['positive', 'negative', 'neutral']
}

clf = AutoFusionClassifier(config)
clf.fit(train_dataframe)           # trains ML backbone, gathers LLM scores, fits fus
pred = clf.predict(["This is amazing!"]) # fused prediction
```

⁷⁰ CLI and Configuration

⁷¹ Users can generate a starter config and train via the command line:

- ⁷² ▪ Create config: `python train_fusion.py --create-config fusion_config.yaml`
- ⁷³ ▪ Train: `python train_fusion.py --config fusion_config.yaml`
- ⁷⁴ ▪ Optional test data and output artifacts are also supported.

75 Quality Control

76 The repository ships legacy unit tests under tests/evaluation/old/ that cover configuration
 77 handling, core types, and package integration. Fusion-specific logic is currently exercised
 78 through CLI-driven workflows and notebooks that run end-to-end training with deterministic
 79 seeds where applicable.

80 Evaluation scripts (tests/evaluation/) provide comprehensive benchmarking on standard
 81 datasets: - **AG News** (Zhang et al., 2015): 4-class topic classification with experiments
 82 across varying training data sizes (20%–100%) - **GoEmotions** (Demszky et al., 2020): 28-class
 83 multi-label emotion classification for validating multi-label fusion performance

84 LLM scoring paths implement retries and disk caching; transformer training supports standard
 85 sanity checks (overfit a small batch, reduced batch sizes for constrained hardware). Metrics
 86 (accuracy/F1, per-label scores) are computed automatically and stored with run artifacts to
 87 facilitate regression tracking and reproducibility.

88 Availability and Installation

89 LabelFusion is distributed as part of the textclassify package under the MIT license and
 90 is available at <https://github.com/DataandAIResearch/LabelFusion>. The fusion components
 91 require Python 3.8+ and common scientific Python dependencies (PyTorch, transformers,
 92 scikit-learn, numpy, pandas, PyYAML). Optional plotting depends on matplotlib/seaborn.
 93 Installation and quick-start snippets are provided in the README and FUSION_README.md.

94 Production-Ready Features

95 Beyond the core fusion methodology, LabelFusion includes features for practical deployment:

- 96 ▪ **LLM Response Caching:** Optional disk-backed caches reuse prior predictions when cache
 97 paths are supplied, with dataset hashes to flag inconsistent inputs.
- 98 ▪ **Results Management:** Built-in ResultsManager tracks experiments, stores predictions,
 99 and computes metrics automatically. Supports comparison across runs and configuration
 100 tracking.
- 101 ▪ **Batch Processing:** Efficient batched scoring of texts with configurable batch sizes for
 102 both ML and LLM components.

103 Impact and Use Cases

104 Empirical Performance

105 LabelFusion has been evaluated on standard benchmark datasets to validate its effectiveness.
 106 Key findings demonstrate consistent improvements over individual model components:

107 AG News Topic Classification

108 Evaluation on the AG News dataset (Zhang et al., 2015) (4-class topic classification) with
 109 5,000 test samples shows:

Training Data	Model	Accuracy	F1-Score	Precision	Recall
20% (800)	Fusion	92.2%	0.922	0.923	0.922
20% (800)	RoBERTa	89.8%	0.899	0.902	0.898
20% (800)	OpenAI	84.4%	0.844	0.857	0.844
40% (1,600)	Fusion	92.2%	0.922	0.924	0.922
40% (1,600)	RoBERTa	91.0%	0.911	0.913	0.910
40% (1,600)	OpenAI	84.4%	0.844	0.857	0.844
100% (4,000)	Fusion	92.4%	0.924	0.926	0.924

Training Data	Model	Accuracy	F1-Score	Precision	Recall
100% (4,000)	RoBERTa	92.2%	0.922	0.923	0.922
100% (4,000)	OpenAI	84.4%	0.844	0.857	0.844

¹¹⁰ **Key Observations:** - Fusion consistently outperforms individual models across all training
¹¹¹ data sizes - With only 20% training data, Fusion achieves 92.2% accuracy—matching its
¹¹² performance with full data - Demonstrates superior **data efficiency**: fusion learning extracts
¹¹³ maximum value from limited examples - RoBERTa alone requires 100% of data to approach
¹¹⁴ Fusion's 20% performance - LLM (OpenAI) shows stable but lower performance, highlighting
¹¹⁵ the value of combining approaches

¹¹⁶ These results validate that learned fusion captures complementary strengths: the LLM provides
¹¹⁷ robust reasoning even with limited training data, while the ML backbone adds efficiency and
¹¹⁸ domain-specific patterns.

¹¹⁹ Application Domains

¹²⁰ Learned fusion excels in scenarios where model strengths complement each other:

- ¹²¹ ▪ **Customer feedback analysis** with nuanced multi-label taxonomies where LLMs handle
¹²² ambiguous sentiment while ML models efficiently process clear cases
- ¹²³ ▪ **Content moderation** where uncertain cases benefit from LLM reasoning while routine
¹²⁴ items rely on the fast ML backbone, enabling real-time processing with accuracy
¹²⁵ guarantees
- ¹²⁶ ▪ **Scientific literature classification** across heterogeneous topics where domain shift is
¹²⁷ common and LLMs provide robustness to new terminology
- ¹²⁸ ▪ **Low-resource settings** where limited training data is available but task complexity requires
¹²⁹ sophisticated reasoning

¹³⁰ The approach enables pragmatic cost control (e.g., the fusion layer learns when to rely more
¹³¹ heavily on the efficient ML backbone versus the more expensive LLM signal) while retaining a
¹³² single trainable decision surface that optimizes for the specific deployment constraints.

¹³³ Acknowledgements

¹³⁴ We thank contributors and users who reported issues and shared datasets. LabelFusion builds on
¹³⁵ the open-source ecosystem, notably Hugging Face Transformers ([Wolf et al., 2019](#)), scikit-learn
¹³⁶ ([Pedregosa et al., 2011](#)), PyTorch ([Paszke et al., 2019](#)), and LLM provider SDKs. The work
¹³⁷ presented in this paper was conducted independently by the author Melchizedek Mashiku and
¹³⁸ is not affiliated with Tanaq Management Services LLC, Contracting Agency to the Division
¹³⁹ of Viral Diseases, Centers for Disease Control and Prevention, Chamblee, Georgia, USA. We
¹⁴⁰ acknowledge the use of the AG News and GoEmotions benchmark datasets for evaluation.

¹⁴¹ Appendix A: Task Formalization

¹⁴² Formally, multi-class classification assigns each input $x \in \mathcal{X}$ to exactly one label among K
¹⁴³ mutually exclusive classes:

$$f_{\text{mc}} : \mathcal{X} \rightarrow \{1, \dots, K\}.$$

¹⁴⁴ In contrast, multi-label classification predicts a subset of relevant classes, represented as a
¹⁴⁵ binary indicator vector $y \in \{0, 1\}^K$, where $y_k = 1$ denotes membership in class k :

$$f_{\text{ml}} : \mathcal{X} \rightarrow \{0, 1\}^K.$$

¹⁴⁶ This distinction shapes the training and inference stack. Multi-class models typically pair
¹⁴⁷ a softmax activation with categorical cross-entropy, yielding normalized class probabilities

148 (Goodfellow et al., 2016). Multi-label classifiers instead apply independent sigmoid activations
 149 with binary cross-entropy, producing class-wise confidence scores that require calibrated
 150 thresholds at prediction time (Goodfellow et al., 2016). LabelFusion preserves these per-class
 151 semantics when concatenating transformer logits and LLM scores, allowing the fusion network
 152 to learn how much to trust each source under either regime.

153 References

- 154 Demszky, D., Movshovitz-Attias, D., Ko, J., Cowen, A., Nemade, G., & Ravi, S. (2020).
 155 GoEmotions: A dataset of fine-grained emotions. *Proceedings of the 58th Annual Meeting*
 156 *of the Association for Computational Linguistics*, 4040–4054. <https://doi.org/10.18653/v1/2020.acl-main.372>
- 158 Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of deep
 159 bidirectional transformers for language understanding. *arXiv Preprint arXiv:1810.04805*.
 160 <https://doi.org/10.48550/arXiv.1810.04805>
- 161 Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press. <https://doi.org/10.7551/mitpress/10251.001.0001>
- 163 Kant, G., Wiebelt, L., Weisser, C., Kis-Katos, K., Luber, M., & Säfken, B. (2022). An
 164 iterative topic model filtering framework for short and noisy user-generated data: Analyzing
 165 conspiracy theories on twitter. *International Journal of Data Science and Analytics*, 20(2),
 166 269–289. <https://doi.org/10.1007/s41060-022-00321-4>
- 167 Kant, G., Zhelyazkov, I., Thielmann, A., Weisser, C., Schlee, M., Ehrling, C., Säfken, B., &
 168 Kneib, T. (2024). One-way ticket to the moon? An NLP-based insight on the phenomenon
 169 of small-scale neo-broker trading. *Social Network Analysis and Mining*, 14(1), 121. <https://doi.org/10.1007/s13278-024-01273-2>
- 171 Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer,
 172 L., & Stoyanov, V. (2019). RoBERTa: A robustly optimized BERT pretraining approach.
 173 *arXiv Preprint arXiv:1907.11692*. <https://doi.org/10.48550/arXiv.1907.11692>
- 174 Luber, M., Weisser, C., Säfken, B., Silbersdorff, A., Kneib, T., & Kis-Katos, K. (2021).
 175 Identifying topical shifts in twitter streams: An integration of non-negative matrix factoriza-
 176 tion, sentiment analysis and structural break models for large scale data. In J. Bright, A.
 177 Giachanou, V. Spaiser, F. Spezzano, A. George, & A. Pavliuc (Eds.), *Disinformation in open*
 178 *online media* (pp. 33–49). Springer International Publishing. ISBN: 978-3-030-87031-7
- 179 OpenAI. (2023). GPT-4 technical report. *arXiv Preprint arXiv:2303.08774*. <https://doi.org/10.48550/arXiv.2303.08774>
- 181 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin,
 182 Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-
 183 performance deep learning library. *Advances in Neural Information Processing Systems*, 32.
 184 <https://doi.org/10.48550/arXiv.1912.01703>
- 185 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel,
 186 M., Prettenhofer, P., Weiss, R., Dubourg, V., & others. (2011). Scikit-learn: Machine
 187 learning in python. *Journal of Machine Learning Research*, 12(Oct), 2825–2830. <http://www.jmlr.org/papers/v12/pedregosa11a.html>
- 189 Thielmann, A. F., Weisser, C., & Säfken, B. (2024). Human in the loop: How to effectively
 190 create coherent topics by manually labeling only a few documents per class. In N. Calzolari,
 191 M.-Y. Kan, V. Hoste, A. Lenci, S. Sakti, & N. Xue (Eds.), *Proceedings of the 2024 joint*
 192 *international conference on computational linguistics, language resources and evaluation*
 193 (*LREC-COLING 2024*) (pp. 8395–8405). ELRA; ICCL. <https://doi.org/10.18653/v1/2024.lrec-main.736>

- 195 Thielmann, A., Weisser, C., & Krenz, A. (2021). One-class support vector machine and LDA
196 topic model integration—evidence for AI patents. In N. H. Phuong & V. Kreinovich (Eds.),
197 *Soft computing: Biomedical and related applications* (pp. 263–272). Springer International
198 Publishing. https://doi.org/10.1007/978-3-030-76620-7_23
- 199 Thielmann, A., Weisser, C., Krenz, A., & Säfken, B. (2021). Unsupervised document
200 classification integrating web scraping, one-class SVM and LDA topic modelling. *Journal
201 of Applied Statistics*, 50(3), 574–591. <https://doi.org/10.1080/02664763.2021.1919063>
- 202 Thormann, M.-L., Farchmin, J., Weisser, C., Kruse, R.-M., Säfken, B., & Silbersdorff, A.
203 (2021). Stock price predictions with LSTM neural networks and twitter sentiment. *Statistics,
204 Optimization & Information Computing*, 9(2), 268–287. [https://doi.org/10.19139/
205 soic-2310-5070-1202](https://doi.org/10.19139/soic-2310-5070-1202)
- 206 Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T.,
207 Louf, R., Funtowicz, M., & others. (2019). HuggingFace’s transformers: State-of-the-art
208 natural language processing. *arXiv Preprint arXiv:1910.03771*. [https://doi.org/10.48550/
209 arXiv.1910.03771](https://doi.org/10.48550/arXiv.1910.03771)
- 210 Zhang, X., Zhao, J., & LeCun, Y. (2015). Character-level convolutional networks for text
211 classification. *Advances in Neural Information Processing Systems*, 28, 649–657. [https:
212 //doi.org/10.48550/arXiv.1509.01626](https://doi.org/10.48550/arXiv.1509.01626)

DRAFT