

¹ LabelFusion: Learning to Fuse LLMs and Transformer Classifiers for Robust Text Classification

³ Michael Schlee¹, Christoph Weisser¹, Timo Kivimäki², Melchizedek
⁴ Mashiku⁴, and Benjamin Saefken³

⁵ 1 Centre for Statistics, Georg-August-Universität Göttingen, Germany ² Department of Politics and
⁶ International Studies, University of Bath, Bath, UK ³ Institute of Mathematics, Clausthal University of
⁷ Technology, Clausthal-Zellerfeld, Germany ⁴ Tanaq Management Services LLC, Contracting Agency to
⁸ the Division of Viral Diseases, Centers for Disease Control and Prevention, Chamblee, Georgia, USA

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).
21
22
23

⁹ Summary

¹⁰ LabelFusion is a fusion ensemble for text classification that learns to combine a traditional
¹¹ transformer-based classifier (e.g., RoBERTa) with one or more Large Language Models (LLMs)
¹² such as OpenAI GPT, Google Gemini, or DeepSeek to deliver accurate and cost-aware predictions
¹³ across multi-class and multi-label tasks. The package provides a simple high-level interface
¹⁴ (AutoFusionClassifier) that trains the full pipeline end-to-end with minimal configuration,
¹⁵ and a flexible API for advanced users. Under the hood, LabelFusion concatenates vector signals
¹⁶ from the ML backbone (logits) and LLM(s) (per-class scores) and trains a compact multi-layer
¹⁷ perceptron (FusionMLP) to produce the final prediction. This learned fusion approach captures
¹⁸ complementary strengths of LLM reasoning and traditional transformer-based classifiers,
¹⁹ yielding robust performance across domains—achieving 92.4% accuracy on AG News topic
²⁰ classification—while enabling practical trade-offs between accuracy, latency, and cost.

²¹ Statement of Need

²² Modern text classification spans diverse scenarios, from sentiment analysis (Kant et al., 2024;
Luber et al., 2021; Thormann et al., 2021) to complex topic tagging (Kant et al., 2022; A.
Thielmann, Weisser, Krenz, & Säfken, 2021; A. Thielmann, Weisser, & Krenz, 2021; A. F.
Thielmann et al., 2024), often under constraints that vary per deployment (throughput, cost
ceilings, data privacy). While transformer classifiers such as BERT/RoBERTa achieve strong
supervised performance (Devlin et al., 2018; Liu et al., 2019), frontier LLMs can excel in
low-data, ambiguous, or cross-domain settings (OpenAI, 2023). No single model family is
typically uniformly best: LLMs are powerful, but comparatively costly, whereas fine-tuned
transformers are efficient but may struggle with out-of-distribution cases.

²³ LabelFusion addresses this gap by: (1) exposing a minimal “AutoFusion” interface that trains a
²⁴ learned combination of an ML backbone and one or more LLMs; (2) supporting both multi-class
²⁵ and multi-label classification; (3) providing a lightweight fusion learner that directly fits on LLM
²⁶ scores and ML logits; and (4) integrating cleanly with existing ensemble utilities. Researchers
²⁷ and practitioners can therefore leverage LLMs where they add value while retaining the speed
²⁸ and determinism of transformer models.

²⁹ State of the Field

³⁰ In applied NLP, common tools such as scikit-learn (Pedregosa et al., 2011) and Hugging Face
³¹ Transformers (Wolf et al., 2019) offer strong baselines but do not provide a learned fusion of
³² LLMs with supervised transformers. Orchestration frameworks (e.g., LangChain) focus on tool
³³ use rather than classification ensembles. LabelFusion contributes a focused, production-minded

⁴² implementation of a small learned combiner that operates on per-class signals from both model
⁴³ families.

⁴⁴ Functionality and Design

⁴⁵ LabelFusion consists of three layers:

- ⁴⁶ ▪ ML component: a RoBERTa-style classifier produces per-class logits for input texts.
- ⁴⁷ ▪ LLM component(s): provider-specific classifiers (OpenAI, Gemini, DeepSeek) return
- ⁴⁸ per-class scores via prompting. Scores can be cached to minimize API calls when cache
- ⁴⁹ locations are provided.
- ⁵⁰ ▪ Fusion component: a compact MLP concatenates ML logits and LLM scores and outputs
- ⁵¹ fused logits. The ML backbone is trained/fine-tuned with a small learning rate; the fusion
- ⁵² MLP uses a higher rate, enabling rapid adaptation without destabilizing the encoder.

⁵³ Key features:

- ⁵⁴ ▪ **Multi-class and multi-label support** with consistent data structures and unified training
- ⁵⁵ pipeline.
- ⁵⁶ ▪ **Optional LLM response caching** reuses on-disk predictions when cache paths are supplied,
- ⁵⁷ with dataset-hash validation to guard against stale files.
- ⁵⁸ ▪ **Batched scoring** processes multiple texts efficiently with configurable batch sizes for both
- ⁵⁹ ML tokenization and LLM API calls.
- ⁶⁰ ▪ **Results management** via ResultsManager tracks experiments, stores predictions, com-
- ⁶¹ putes metrics, and enables reproducible research workflows.
- ⁶² ▪ **Flexible interfaces**: Command-line training via `train_fusion.py` with YAML configs for
- ⁶³ research; or minimal AutoFusion API for quick deployment.
- ⁶⁴ ▪ **Composable design**: LabelFusion can serve as a strong base learner in higher-level
- ⁶⁵ ensembles (e.g., voting/weighted combinations of multiple fusion models).

⁶⁶ We support both multi-class setups (one label per input) and multi-label scenarios (multi-

⁶⁷ ple labels per input), and point readers to Appendix A for formal definitions and training

⁶⁸ implications.

⁶⁹ Minimal Example (AutoFusion)

```
from textclassify import AutoFusionClassifier

# Multi-class: exactly one of the sentiment labels applies
multiclass_config = {
    'llm_provider': 'deepseek',
    'label_columns': ['positive', 'negative', 'neutral'],
    'multi_label': False
}
multiclass_clf = AutoFusionClassifier(multiclass_config)
multiclass_clf.fit(train_dataframe)
multiclass_pred = multiclass_clf.predict(["This is amazing!"])

# Multi-label: news article can belong to several topics simultaneously
multilabel_config = {
    'llm_provider': 'deepseek',
    'label_columns': ['politics', 'economy', 'technology'],
    'multi_label': True
}
multilabel_clf = AutoFusionClassifier(multilabel_config)
multilabel_clf.fit(train_dataframe)
multilabel_pred = multilabel_clf.predict(["New investment in AI chips"])
```

70 **CLI and Configuration**

71 Users can generate a starter config and train via the command line:

- 72 ■ Create config: `python train_fusion.py --create-config fusion_config.yaml`
- 73 ■ Train: `python train_fusion.py --config fusion_config.yaml`
- 74 ■ Optional test data and output artifacts are also supported.

75 **Quality Control**

76 The repository ships legacy unit tests under `tests/evaluation/old/` that cover configuration
77 handling, core types, and package integration. Fusion-specific logic is currently exercised
78 through CLI-driven workflows and notebooks that run end-to-end training with deterministic
79 seeds where applicable.

80 Evaluation scripts (`tests/evaluation/`) provide comprehensive benchmarking on standard
81 datasets: - **AG News** ([Zhang et al., 2015](#)): 4-class topic classification with experiments
82 across varying training data sizes (20%–100%) - **GoEmotions** ([Demszky et al., 2020](#)): 28-class
83 multi-label emotion classification for validating multi-label fusion performance

84 LLM scoring paths implement retries and disk caching; transformer training supports standard
85 sanity checks (overfit a small batch, reduced batch sizes for constrained hardware). Metrics
86 (accuracy/F1, per-label scores) are computed automatically and stored with run artifacts to
87 facilitate regression tracking and reproducibility.

88 **Availability and Installation**

89 LabelFusion is distributed as part of the `textclassify` package under the MIT license and
90 is available at <https://github.com/DataandAIResearch/LabelFusion>. The fusion components
91 require Python 3.8+ and common scientific Python dependencies (PyTorch, transformers,
92 scikit-learn, numpy, pandas, PyYAML). Optional plotting depends on matplotlib/seaborn.
93 Installation and quick-start snippets are provided in the `README` and `FUSION_README.md`.

94 **Production-Ready Features**

95 Beyond the core fusion methodology, LabelFusion includes features for practical deployment:

- 96 ■ **LLM Response Caching:** Optional disk-backed caches reuse prior predictions when cache
97 paths are supplied, with dataset hashes to flag inconsistent inputs.
- 98 ■ **Results Management:** Built-in `ResultsManager` tracks experiments, stores predictions,
99 and computes metrics automatically. Supports comparison across runs and configuration
100 tracking.
- 101 ■ **Batch Processing:** Efficient batched scoring of texts with configurable batch sizes for
102 both ML and LLM components.

103 **Impact and Use Cases**

104 **Empirical Performance**

105 LabelFusion has been evaluated on standard benchmark datasets to validate its effectiveness.
106 Key findings demonstrate consistent improvements over individual model components:

107 **AG News Topic Classification**

108 Evaluation on the AG News dataset ([Zhang et al., 2015](#)) (4-class topic classification) with
109 5,000 test samples shows:

Training Data	Model	Accuracy	F1-Score	Precision	Recall
20% (800)	Fusion	92.2%	0.922	0.923	0.922
20% (800)	RoBERTa	89.8%	0.899	0.902	0.898
20% (800)	OpenAI	84.4%	0.844	0.857	0.844
40% (1,600)	Fusion	92.2%	0.922	0.924	0.922
40% (1,600)	RoBERTa	91.0%	0.911	0.913	0.910
40% (1,600)	OpenAI	84.4%	0.844	0.857	0.844
100% (4,000)	Fusion	92.4%	0.924	0.926	0.924
100% (4,000)	RoBERTa	92.2%	0.922	0.923	0.922
100% (4,000)	OpenAI	84.4%	0.844	0.857	0.844

- ¹¹⁰ **Key Observations:** - Fusion consistently outperforms individual models across all training
¹¹¹ data sizes - With only 20% training data, Fusion achieves 92.2% accuracy—matching its
¹¹² performance with full data - Demonstrates superior **data efficiency**: fusion learning extracts
¹¹³ maximum value from limited examples - RoBERTa alone requires 100% of data to approach
¹¹⁴ Fusion's 20% performance - LLM (OpenAI) shows stable but lower performance, highlighting
¹¹⁵ the value of combining approaches
- ¹¹⁶ These results validate that learned fusion captures complementary strengths: the LLM provides
¹¹⁷ robust reasoning even with limited training data, while the ML backbone adds efficiency and
¹¹⁸ domain-specific patterns.

¹¹⁹ Application Domains

¹²⁰ Learned fusion excels in scenarios where model strengths complement each other:

- ¹²¹ ▪ **Customer feedback analysis** with nuanced multi-label taxonomies where LLMs handle
¹²² ambiguous sentiment while ML models efficiently process clear cases
- ¹²³ ▪ **Content moderation** where uncertain cases benefit from LLM reasoning while rou-
¹²⁴ tine items rely on the fast ML backbone, enabling real-time processing with accuracy
¹²⁵ guarantees
- ¹²⁶ ▪ **Scientific literature classification** across heterogeneous topics where domain shift is
¹²⁷ common and LLMs provide robustness to new terminology
- ¹²⁸ ▪ **Low-resource settings** where limited training data is available but task complexity requires
¹²⁹ sophisticated reasoning

¹³⁰ The approach enables pragmatic cost control (e.g., the fusion layer learns when to rely more
¹³¹ heavily on the efficient ML backbone versus the more expensive LLM signal) while retaining a
¹³² single trainable decision surface that optimizes for the specific deployment constraints.

¹³³ Acknowledgements

¹³⁴ We thank contributors and users who reported issues and shared datasets. LabelFusion builds on
¹³⁵ the open-source ecosystem, notably Hugging Face Transformers ([Wolf et al., 2019](#)), scikit-learn
¹³⁶ ([Pedregosa et al., 2011](#)), PyTorch ([Paszke et al., 2019](#)), and LLM provider SDKs. The work
¹³⁷ presented in this paper was conducted independently by the author Melchizedek Mashiku and
¹³⁸ is not affiliated with Tanaq Management Services LLC, Contracting Agency to the Division
¹³⁹ of Viral Diseases, Centers for Disease Control and Prevention, Chamblee, Georgia, USA. We
¹⁴⁰ acknowledge the use of the AG News and GoEmotions benchmark datasets for evaluation.

¹⁴¹ Appendix A: Task Formalization

¹⁴² Formally, multi-class classification assigns each input $x \in \mathcal{X}$ to exactly one label among K
¹⁴³ mutually exclusive classes:

$$f_{mc} : \mathcal{X} \rightarrow \{1, \dots, K\}.$$

¹⁴⁴ In contrast, multi-label classification predicts a subset of relevant classes, represented as a
¹⁴⁵ binary indicator vector $\mathbf{y} \in \{0, 1\}^K$, where $y_k = 1$ denotes membership in class k :

$$f_{\text{ml}} : \mathcal{X} \rightarrow \{0, 1\}^K.$$

¹⁴⁶ This distinction shapes the training and inference stack. Multi-class models typically pair
¹⁴⁷ a softmax activation with categorical cross-entropy, yielding normalized class probabilities
¹⁴⁸ (Goodfellow et al., 2016). Multi-label classifiers instead apply independent sigmoid activations
¹⁴⁹ with binary cross-entropy, producing class-wise confidence scores that require calibrated
¹⁵⁰ thresholds at prediction time (Goodfellow et al., 2016). LabelFusion preserves these per-class
¹⁵¹ semantics when concatenating transformer logits and LLM scores, allowing the fusion network
¹⁵² to learn how much to trust each source under either regime.

153 References

- ¹⁵⁴ Demszky, D., Movshovitz-Attias, D., Ko, J., Cowen, A., Nemade, G., & Ravi, S. (2020).
¹⁵⁵ GoEmotions: A dataset of fine-grained emotions. *Proceedings of the 58th Annual Meeting*
¹⁵⁶ *of the Association for Computational Linguistics*, 4040–4054. <https://doi.org/10.18653/v1/2020.acl-main.372>
- ¹⁵⁷ Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2018). BERT: Pre-training of deep
¹⁵⁸ bidirectional transformers for language understanding. *arXiv Preprint arXiv:1810.04805*.
¹⁵⁹ <https://doi.org/10.48550/arXiv.1810.04805>
- ¹⁶⁰ Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press. <https://doi.org/10.7551/mitpress/10251.001.0001>
- ¹⁶¹ Kant, G., Wiebelt, L., Weisser, C., Kis-Katos, K., Luber, M., & Säfken, B. (2022). An
¹⁶² iterative topic model filtering framework for short and noisy user-generated data: Analyzing
¹⁶³ conspiracy theories on twitter. *International Journal of Data Science and Analytics*, 20(2),
¹⁶⁴ 269–289. <https://doi.org/10.1007/s41060-022-00321-4>
- ¹⁶⁵ Kant, G., Zhelyazkov, I., Thielmann, A., Weisser, C., Schlee, M., Ehrling, C., Säfken, B., &
¹⁶⁶ Kneib, T. (2024). One-way ticket to the moon? An NLP-based insight on the phenomenon
¹⁶⁷ of small-scale neo-broker trading. *Social Network Analysis and Mining*, 14(1), 121. <https://doi.org/10.1007/s13278-024-01273-2>
- ¹⁶⁸ Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer,
¹⁶⁹ L., & Stoyanov, V. (2019). RoBERTa: A robustly optimized BERT pretraining approach.
¹⁷⁰ *arXiv Preprint arXiv:1907.11692*. <https://doi.org/10.48550/arXiv.1907.11692>
- ¹⁷¹ Luber, M., Weisser, C., Säfken, B., Silbersdorff, A., Kneib, T., & Kis-Katos, K. (2021).
¹⁷² Identifying topical shifts in twitter streams: An integration of non-negative matrix factorisation,
¹⁷³ sentiment analysis and structural break models for large scale data. In J. Bright, A.
¹⁷⁴ Giachanou, V. Spaiser, F. Spezzano, A. George, & A. Pavliuc (Eds.), *Disinformation in open*
¹⁷⁵ *online media* (pp. 33–49). Springer International Publishing. ISBN: 978-3-030-87031-7
- ¹⁷⁶ OpenAI. (2023). GPT-4 technical report. *arXiv Preprint arXiv:2303.08774*. <https://doi.org/10.48550/arXiv.2303.08774>
- ¹⁷⁷ Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin,
¹⁷⁸ Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-
¹⁷⁹ performance deep learning library. *Advances in Neural Information Processing Systems*, 32.
¹⁸⁰ <https://doi.org/10.48550/arXiv.1912.01703>
- ¹⁸¹ Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel,
¹⁸² M., Prettenhofer, P., Weiss, R., Dubourg, V., & others. (2011). Scikit-learn: Machine
¹⁸³ learning in python. *Journal of Machine Learning Research*, 12(Oct), 2825–2830. <http://www.jmlr.org/papers/v12/pedregosa11a.html>

- 189 Thielmann, A. F., Weisser, C., & Säfken, B. (2024). Human in the loop: How to effectively
190 create coherent topics by manually labeling only a few documents per class. In N. Calzolari,
191 M.-Y. Kan, V. Hoste, A. Lenci, S. Sakti, & N. Xue (Eds.), *Proceedings of the 2024 joint*
192 *international conference on computational linguistics, language resources and evaluation*
193 (*LREC-COLING 2024*) (pp. 8395–8405). ELRA; ICCL. <https://doi.org/10.18653/v1/2024.lrec-main.736>
- 195 Thielmann, A., Weisser, C., & Krenz, A. (2021). One-class support vector machine and LDA
196 topic model integration—evidence for AI patents. In N. H. Phuong & V. Kreinovich (Eds.),
197 *Soft computing: Biomedical and related applications* (pp. 263–272). Springer International
198 Publishing. https://doi.org/10.1007/978-3-030-76620-7_23
- 199 Thielmann, A., Weisser, C., Krenz, A., & Säfken, B. (2021). Unsupervised document
200 classification integrating web scraping, one-class SVM and LDA topic modelling. *Journal*
201 *of Applied Statistics*, 50(3), 574–591. <https://doi.org/10.1080/02664763.2021.1919063>
- 202 Thormann, M.-L., Farchmin, J., Weisser, C., Kruse, R.-M., Säfken, B., & Silbersdorff, A.
203 (2021). Stock price predictions with LSTM neural networks and twitter sentiment. *Statistics,*
204 *Optimization & Information Computing*, 9(2), 268–287. <https://doi.org/10.19139/soic-2310-5070-1202>
- 206 Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T.,
207 Louf, R., Funtowicz, M., & others. (2019). HuggingFace's transformers: State-of-the-art
208 natural language processing. *arXiv Preprint arXiv:1910.03771*. <https://doi.org/10.48550/arXiv.1910.03771>
- 210 Zhang, X., Zhao, J., & LeCun, Y. (2015). Character-level convolutional networks for text
211 classification. *Advances in Neural Information Processing Systems*, 28, 649–657. <https://doi.org/10.48550/arXiv.1509.01626>